

Overview

- CASENUM- Case number to identify the claim, a numeric vector
- ATTORNEY Whether the claimant is represented by an attorney (=1 if yes and =2 if no)
- CLMSEX Claimant's gender (=1 if male and =2 if female), a numeric vector
- CLMINSUR Whether or not the driver of the claimant's vehicle was uninsured (=1 if yes, =2 if no)
- SEATBELT Whether or not the claimant was wearing a seatbelt/child restraint (=1 if yes, =2 if no)
- CLMAGE Claimant's age, a numeric vector
- LOSS The claimant's total economic loss (in thousands)

Problem Statement

```
In [76]: from PIL import ImageGrab
ImageGrab.grabclipboard()
```

Out[76]:

Conditional Expectation:

- Now what to predict?- Only two possibilities for the response variable- Either claimant is attorney, or not.
- Given a set of values of explanatory variables , can we predict with 100% certainty that the claimant is represented by an attorney/not?
- Of Course, the only thing we can sensibly try to predict is the probability that the claimant is represented by an attorney given a set of values for the explanatory variables

```
In [ ]: # Importing libraries
```

```
In [11]: import pandas as pd
import numpy as np
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stat
import warnings
warnings.filterwarnings('ignore')
from sklearn.linear_model import LogisticRegression
```

```
In [ ]: # Load the dataset
```

```
In [12]: df = pd.read_csv('claimants.csv')
df.head()
```

```
Out[12]:
```

CASENUM	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
0	5	0	0.0	1.0	0.0	50.0
34.940						

	CASENUM	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
1	3	1	1.0	0.0	0.0	18.0	0.891
2	66	1	0.0	1.0	0.0	5.0	0.330
3	70	0	0.0	1.0	1.0	31.0	0.037
4	96	1	0.0	1.0	0.0	30.0	0.038

In [9]: *# shape of the dataset*
df.shape

Out[9]: (1340, 7)

In [75]: df.CLMINSUR.unique()

Out[75]: array([1., 0.])

In [13]: *# Dropping the case number columns as it has Unique values and it is not contributing to t*
df.drop('CASENUM', axis=1, inplace = True)
df.head()

Out[13]:

	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
0	0	0.0	1.0	0.0	50.0	34.940
1	1	1.0	0.0	0.0	18.0	0.891
2	1	0.0	1.0	0.0	5.0	0.330
3	0	0.0	1.0	1.0	31.0	0.037
4	1	0.0	1.0	0.0	30.0	0.038

Descriptive Statistics

In [36]: df.describe()

Out[36]:

	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
count	1340.000000	1328.000000	1299.000000	1292.000000	1151.000000	1340.000000
mean	0.488806	0.558735	0.907621	0.017028	28.414422	3.806307
std	0.500061	0.496725	0.289671	0.129425	20.304451	10.636903
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	1.000000	0.000000	9.000000	0.400000
50%	0.000000	1.000000	1.000000	0.000000	30.000000	1.069500
75%	1.000000	1.000000	1.000000	0.000000	43.000000	3.781500
max	1.000000	1.000000	1.000000	1.000000	95.000000	173.604000

In [21]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1340 entries, 0 to 1339

total 6 columns):

#	Column	Non-Null	Count	Dtype
0	ATTORNEY	1340	non-null	int64
1	CLMSEX	1328	non-null	float64
2	CLMINSUR	1299	non-null	float64
3	SEATBELT	1292	non-null	float64
4	CLMAGE	1151	non-null	float64
5	LOSS	1340	non-null	float64

dtypes: float64(5), int64(1)
memory usage: 62.9 KB

Count of Duplicated Rows

In [4]: `df.duplicated().sum()`

Out[4]: 26

In [5]: `df[df.duplicated()].shape`

Out[5]: (26, 6)

print the duplicated rows

In [39]: `df[df.duplicated()]`

Out[39]:

	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
77	1	1.0	1.0	0.0	NaN	0.000
162	1	1.0	1.0	0.0	NaN	0.800
439	1	0.0	1.0	0.0	57.0	3.889
697	1	0.0	1.0	0.0	NaN	0.400
773	0	0.0	1.0	0.0	NaN	4.395
779	1	1.0	1.0	0.0	NaN	0.000
788	1	0.0	1.0	0.0	NaN	0.150
834	0	1.0	1.0	0.0	48.0	1.050
866	1	1.0	1.0	0.0	30.0	0.300
942	0	1.0	1.0	0.0	8.0	0.100
970	1	0.0	1.0	0.0	11.0	0.500
982	1	0.0	1.0	0.0	NaN	1.000
990	0	1.0	1.0	0.0	43.0	8.490
1049	1	1.0	1.0	0.0	NaN	0.050
1075	0	0.0	1.0	0.0	7.0	0.640
1120	0	1.0	1.0	0.0	NaN	1.000
1121	1	1.0	1.0	0.0	NaN	0.500
1132	1	1.0	1.0	0.0	NaN	0.050
1152	1	0.0	1.0	0.0	NaN	0.150
1186	1	0.0	1.0	0.0	NaN	0.040
1231	1	0.0	1.0	0.0	NaN	0.100

	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
1234	1	1.0	1.0	0.0	NaN	1.640
1260	1	0.0	1.0	0.0	0.0	0.440
1269	0	1.0	1.0	0.0	NaN	3.500
1304	1	1.0	1.0	0.0	NaN	0.500
1320	1	1.0	1.0	0.0	NaN	0.540


```
In [23]: df.isnull().sum()
```

```
Out[23]: ATTORNEY      0
CLMSEX      12
CLMINSUR    41
SEATBELT    48
CLMAGE     189
LOSS        0
dtype: int64
```

Missing Values

```
In [40]: plt.figure(figsize = (16,9))
plt.xticks(fontsize = 14)
cols= df.columns
colors = ['#000000', '#7FFF00']
sns.heatmap(df[cols].isnull(), cmap = sns.color_palette(colors))
```

```
Out[40]: <AxesSubplot:>
```



Exploratory Data Analysis

values and SalesPrice

Let's plot some diagram for this relationship



In [27]:


```
features_with_na = [features for features in df.columns if df[features].isnull().sum()>1]

for feature in features_with_na:
    if feature!="CLMAGE":
        data=df.copy()

        # Let's make a variable that indicates 1 if the observation was missing or zero
        data[feature]=np.where(data[feature].isnull(),1,0)

        # Let's calculate the mean SalePrice where the information is missing or present
        data.groupby(feature)['CLMAGE'].median().plot.bar()
        plt.ylabel('CLMAGE')
        plt.title(feature)
        plt.show()
```






In [28]:




```
features_with_na = [features for features in df.columns if df[features].isnull().sum()>1]

for feature in features_with_na:
    data=df.copy()

    # Let's make a variable that indicates 1 if the observation was missing or zero
    data[feature]=np.where(data[feature].isnull(),1,0)

    # Let's calculate the mean SalePrice where the information is missing or present
    data.groupby(feature)['LOSS'].median().plot.bar()
    plt.ylabel('LOSS')
    plt.title(feature)
    plt.show()
```






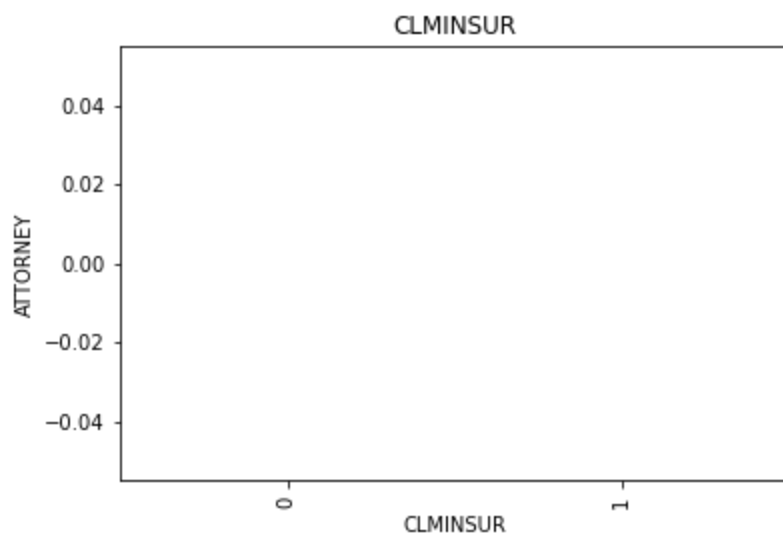
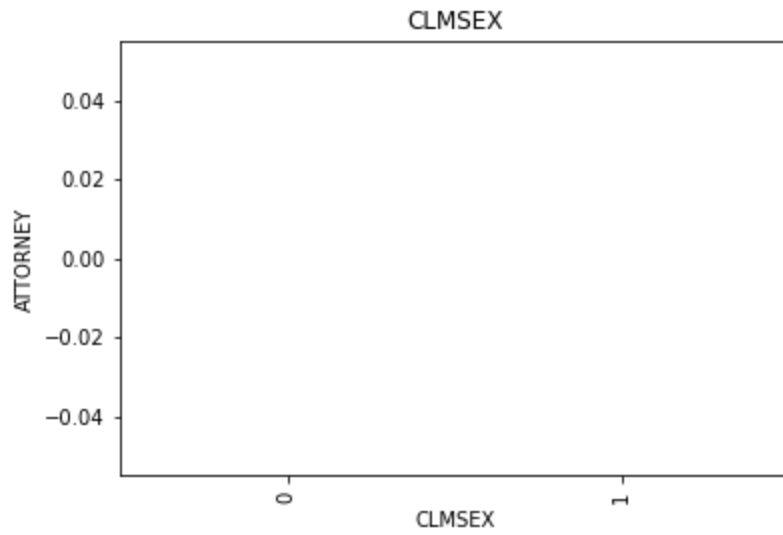
```
In [45]: features_with_na = [features for features in df.columns if df[features].isnull().sum()>1]

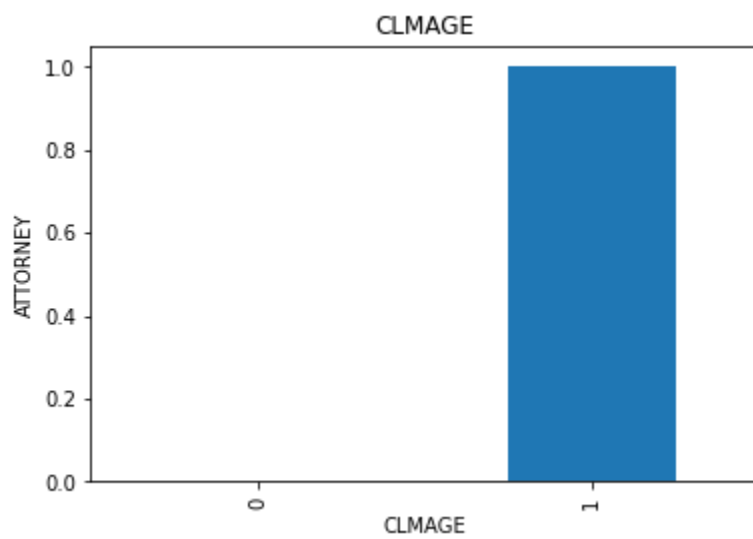
for feature in features_with_na:
    data=df.copy()

    # Let's make a variable that indicates 1 if the observation was missing or zero
    data[feature]=np.where(data[feature].isnull(),1,0)

    # Let's calculate the mean SalePrice where the information is missing or present
    data.groupby(feature)['ATTORNEY'].median().plot.bar()
    plt.ylabel('ATTORNEY')
```

```
plt.title(feature)  
plt.show()
```





```
In [7]: df.CLMINSUR.mode()[0]
```

```
Out[7]: 1.0
```

```
In [26]: df.CLMAGE.median()
```

```
Out[26]: 28.414422241529106
```

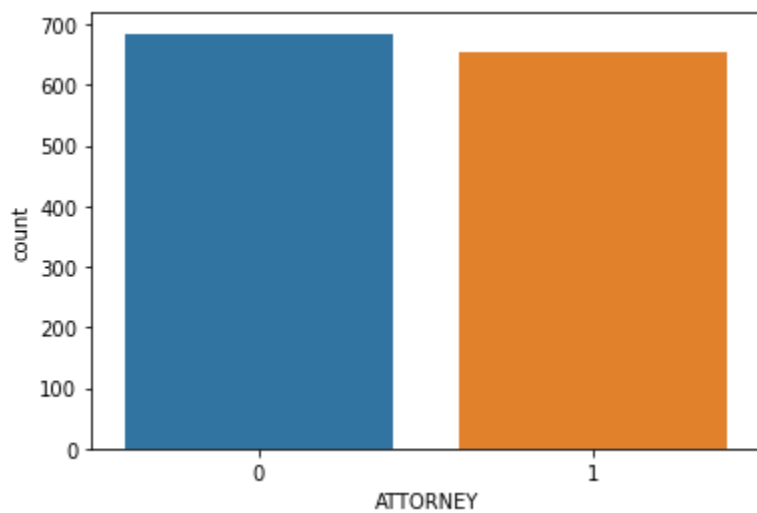
```
In [17]: #Fill nan values with mode of the categorical column
df["CLMSEX"].fillna(df.CLMSEX.mode()[0],inplace=True) # df.CLMSEX.mode() = 1
df["CLMINSUR"].fillna(df.CLMINSUR.mode()[0],inplace=True) # df.CLMINSUR.mode() = 1
df["SEATBELT"].fillna(df.SEATBELT.mode()[0],inplace=True) # df.SEATBELT.mode() = 0

df.CLMAGE.fillna(df.CLMAGE.median(),inplace=True) # df.CLMAGE.median() = 28.41
df.isnull().sum()
```

```
Out[17]: ATTORNEY      0
CLMSEX      0
CLMINSUR    0
SEATBELT    0
CLMAGE      0
LOSS        0
dtype: int64
```


```
In [10]: sns.countplot(x="ATTORNEY",data=df)
```

```
Out[10]: <AxesSubplot:xlabel='ATTORNEY', ylabel='count'>
```



```
In [11]: df['ATTORNEY'].value_counts().plot(kind='pie', autopct='%.1f%%')
```


```
Out[11]: <AxesSubplot:ylabel='ATTORNEY'>
```



The target in the Data is Balanced. One of the value counts is not more than the other approximately equally distributed

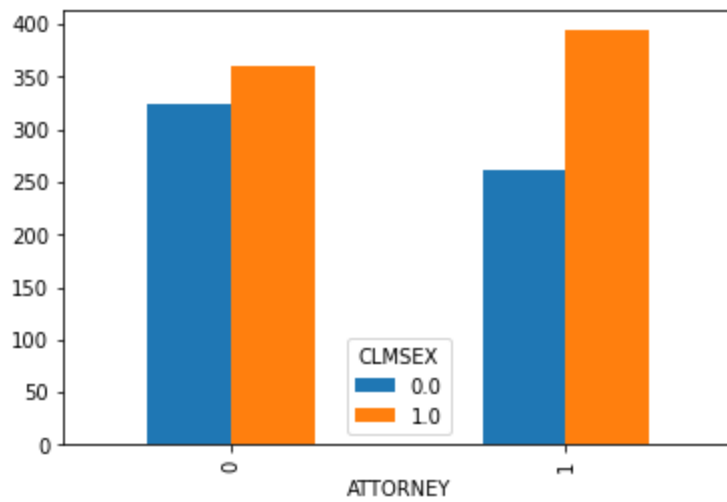
```
In [12]: pd.crosstab(df.ATTORNEY, df.CLMINSUR).plot(kind="bar")
```

```
Out[12]: <AxesSubplot:xlabel='ATTORNEY'>
```



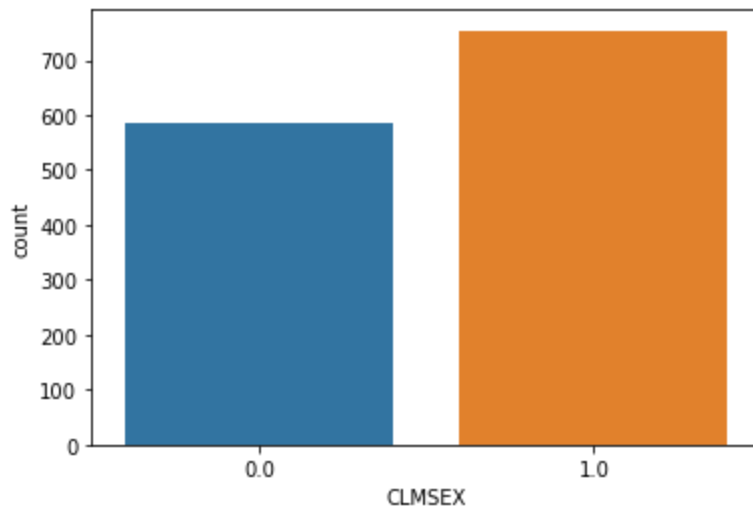
```
In [59]: pd.crosstab(df.ATTORNEY, df.CLMSEX).plot(kind="bar")
```

Out[59]: <AxesSubplot: xlabel='ATTORNEY'>



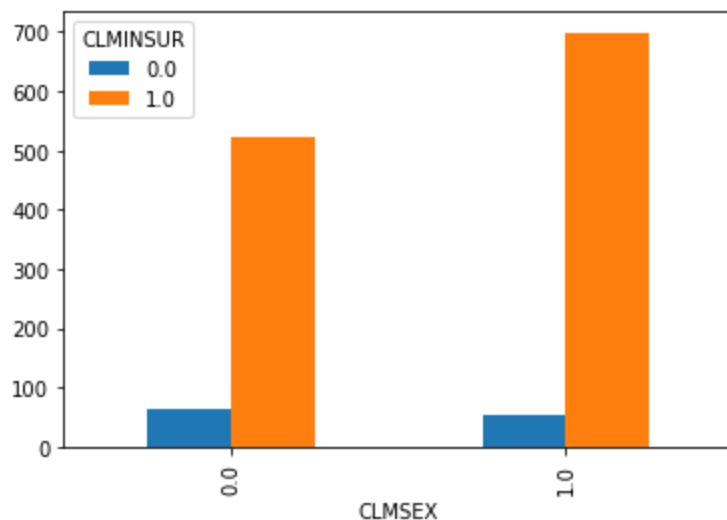
In [13]: `sns.countplot(x="CLMSEX", data=df)`

Out[13]: <AxesSubplot: xlabel='CLMSEX', ylabel='count'>



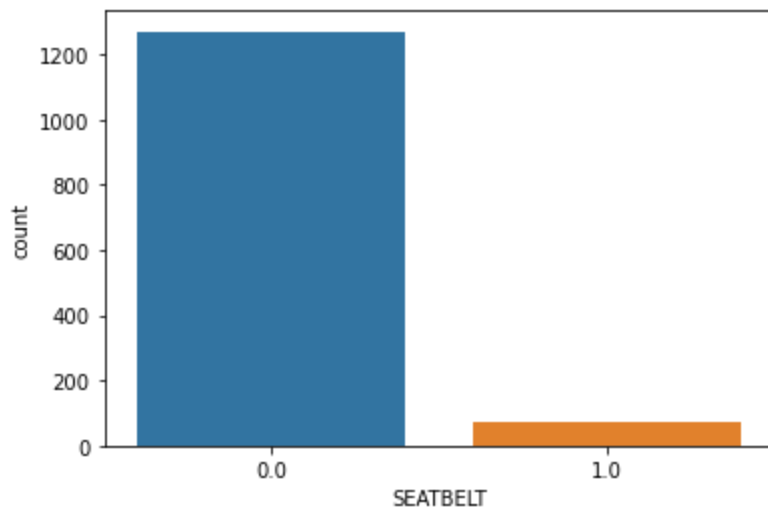
In [14]: `pd.crosstab(df.CLMSEX, df.CLMINSUR).plot(kind="bar")`

Out[14]: <AxesSubplot: xlabel='CLMSEX'>



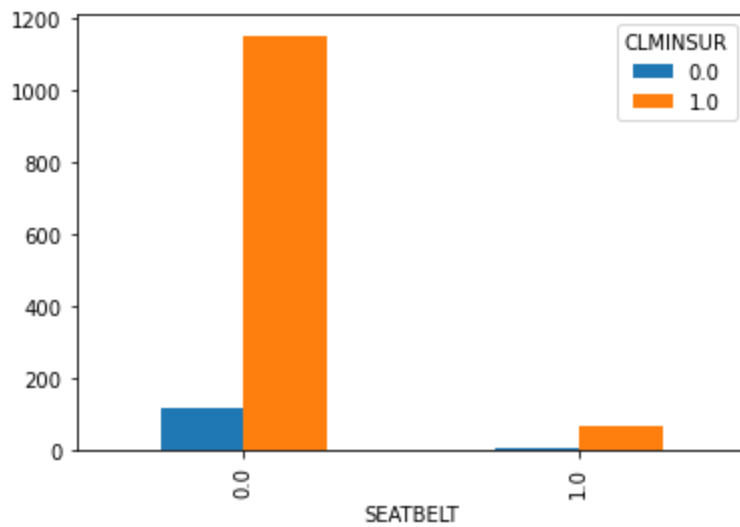
In [15]: `sns.countplot(x="SEATBELT", data=df)`

Out[15]: <AxesSubplot:xlabel='SEATBELT', ylabel='count'>



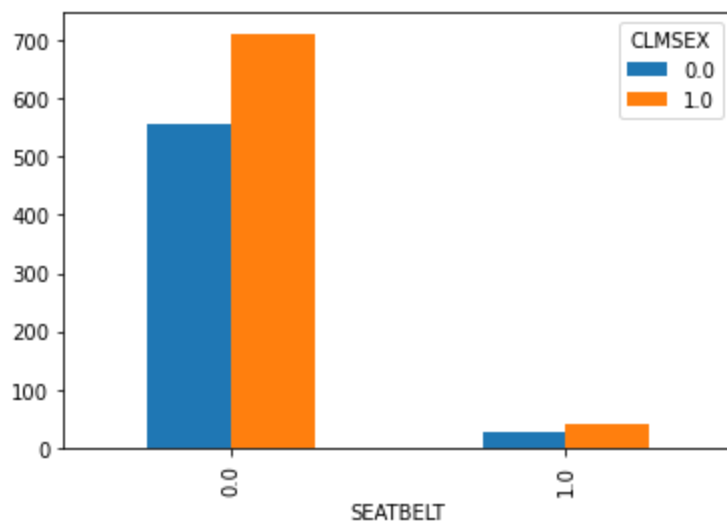
In [16]: `pd.crosstab(df.SEATBELT,df.CLINSUR).plot(kind="bar")`

Out[16]: <AxesSubplot:xlabel='SEATBELT'>



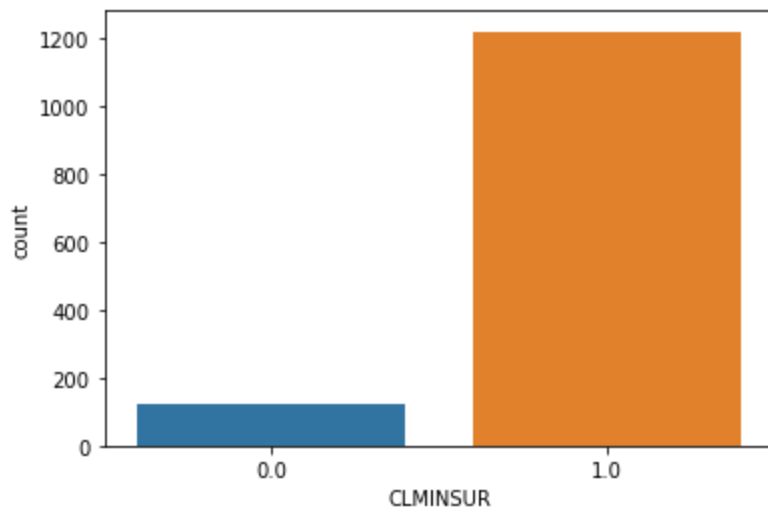
In [58]: `pd.crosstab(df.SEATBELT,df.CLMSEX).plot(kind="bar")`

Out[58]: <AxesSubplot:xlabel='SEATBELT'>




```
sns.countplot(x="CLMINSUR", data=df)
```

Out[17]: <AxesSubplot:xlabel='CLMINSUR', ylabel='count'>



In [18]: `df['CLMSEX'].value_counts().plot(kind='pie', autopct='%0.1f%%')`

Out[18]: <AxesSubplot:ylabel='CLMSEX'>




In [19]: `df.columns`

Out[19]: Index(['ATTORNEY', 'CLMSEX', 'CLMINSUR', 'SEATBELT', 'CLMAGE', 'LOSS'], dtype='object')


In [20]: `df['CLMINSUR'].value_counts().plot(kind='pie', autopct='%0.1f%%')`

Out[20]: <AxesSubplot:ylabel='CLMINSUR'>




```
In [21]: df['SEATBELT'].value_counts().plot(kind='pie', autopct='%0.1f%%')
```

```
Out[21]: <AxesSubplot:ylabel='SEATBELT'>
```




```
In [22]: df['CLMINSUR'].value_counts().plot(kind='pie', autopct='%0.1f%%')
```

```
Out[22]: <AxesSubplot:ylabel='CLMINSUR'>
```




```
In [23]: df['LOSS'].value_counts().plot(kind='hist')
```

```
Out[23]: <AxesSubplot:ylabel='Frequency'>
```



```
In [27]: df['CLMAGE'].value_counts().plot(kind='hist')
```

```
Out[27]: <AxesSubplot:ylabel='Frequency'>
```



```
In [33]: df.CLMAGE.sort_values()
```

```
Out[33]: 1260    0.0
608     0.0
74      0.0
615     0.0
1252    0.0
...
618     83.0
853     84.0
1057    86.0
737     88.0
635     95.0
Name: CLMAGE, Length: 1340, dtype: float64
```

```
In [47]: df[df.CLMAGE.values<18]
```

```
Out[47]:
```

	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
2	1	0.0	1.0	0.0	5.0	0.330
6	0	0.0	1.0	0.0	9.0	3.538
12	1	0.0	1.0	0.0	7.0	1.678
15	1	1.0	0.0	0.0	9.0	0.053

	ATTORNEY	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
18	1	0.0	1.0	0.0	3.0	0.000
...
1328	0	0.0	1.0	0.0	14.0	0.400
1331	0	1.0	1.0	0.0	3.0	0.950
1332	1	1.0	1.0	0.0	9.0	0.000
1334	1	1.0	1.0	0.0	16.0	0.060
1338	0	1.0	0.0	0.0	8.0	3.177

493 rows × 6 columns

```
In [44]: df['CLMAGE'] < 18
```

```
Out[44]: 0      False
1      False
2       True
3      False
4      False
...
1335    False
1336    False
1337    False
1338     True
1339    False
Name: CLMAGE, Length: 1340, dtype: bool
```


```
In [53]: df.columns
```

```
Out[53]: Index(['ATTORNEY', 'CLMSEX', 'CLMINSUR', 'SEATBELT', 'CLMAGE', 'LOSS'], dtype='object')
```

ATTORNEY Yes or No Vs CLMSEX

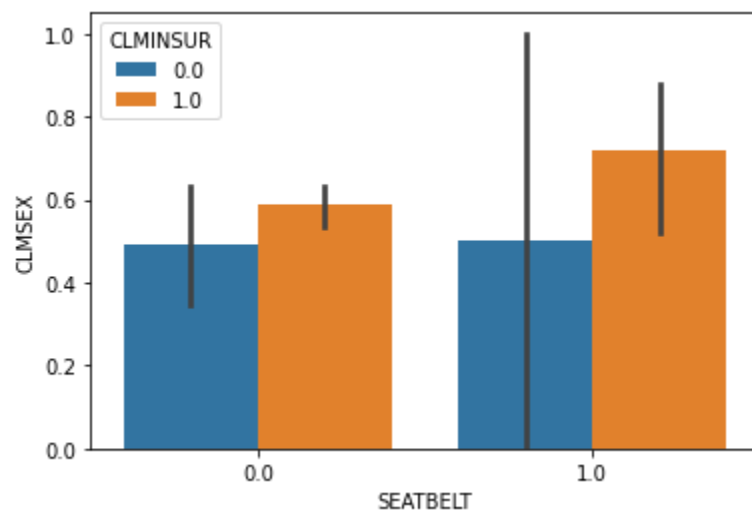
```
In [60]: sns.barplot('ATTORNEY', 'CLMINSUR', hue='CLMSEX', data=df)
```

```
Out[60]: <AxesSubplot: xlabel='ATTORNEY', ylabel='CLMINSUR'>
```




```
In [66]: sns.barplot('SEATBELT', 'CLMSEX', hue='CLMINSUR', data=df[df['CLMAGE'] < 18])
```


Out[66]: <AxesSubplot:xlabel='SEATBELT', ylabel='CLMSEX'>




In [67]: `sns.barplot('SEATBELT', 'CLMSEX', hue='CLMINSUR', data=df[df['CLMAGE']>18])`

Out[67]: <AxesSubplot:xlabel='SEATBELT', ylabel='CLMSEX'>




In [68]: `sns.pairplot(df, hue = 'ATTORNEY')`

Out[68]: <seaborn.axisgrid.PairGrid at 0x219a5bd3e50>




```
In [69]: sns.FacetGrid(df, hue = 'ATTORNEY', size = 6).map(plt.scatter,"LOSS",'CLMAGE').add_legend(plt.show())
```



Outlier Detection using Boxplot


In [71]: `sns.boxplot(data =df,orient = "v")`

Out[71]: <AxesSubplot:>




In [73]: `sns.boxplot(x="ATTORNEY",y="CLIMAGE",data=df,palette = "hls")`

Out[73]: <AxesSubplot:xlabel='ATTORNEY', ylabel='CLIMAGE'>




In [74]: `sns.boxplot(x="ATTORNEY",y="LOSS",data=df,palette="hls")`

Out[74]: `<AxesSubplot:xlabel='ATTORNEY', ylabel='LOSS'>`



In [77]: `sns.boxplot(x="CLMSEX",y="CLMAGE",data=df,palette="hls")`

Out[77]: `<AxesSubplot:xlabel='CLMSEX', ylabel='CLMAGE'>`



In [18]: `#Dividing our data into input and output variables
x = df.iloc[:,1:]
y = df.iloc[:,0]`

In [23]:

```
x
```

Out[23]:

	CLMSEX	CLMINSUR	SEATBELT	CLMAGE	LOSS
0	0.0	1.0	0.0	50.0	34.940
1	1.0	0.0	0.0	18.0	0.891
2	0.0	1.0	0.0	5.0	0.330
3	0.0	1.0	1.0	31.0	0.037
4	0.0	1.0	0.0	30.0	0.038
...
1335	0.0	1.0	0.0	30.0	0.576
1336	1.0	1.0	0.0	46.0	3.705
1337	1.0	1.0	0.0	39.0	0.099
1338	1.0	0.0	0.0	8.0	3.177
1339	1.0	1.0	0.0	30.0	0.688

1340 rows × 5 columns

In [19]:

```
y
```

Out[19]:

```
0      0
1      1
2      1
3      0
4      1
..
1335   1
1336   0
1337   1
1338   0
1339   1
Name: ATTORNEY, Length: 1340, dtype: int64
```

In [20]:

```
# Building a Logistic Regression and fitting the values
classifier = LogisticRegression()
classifier.fit(x,y)
```

Out[20]:

```
LogisticRegression()
```

In [21]:

```
classifier.intercept_
```

Out[21]:

```
array([-0.13616435])
```

In [22]:

```
classifier.coef_
```

Out[22]:

```
array([[ 0.31756644,  0.50430419, -0.52798522,  0.00680276, -0.32257734]])
```

In [83]:

```
# Predicting for x dataset
y_pred = classifier.predict(x)
```

```
In [84]: # Creating a dataframe to with actual value and predicted value
predict = pd.DataFrame({'Actual':y, 'Predicted':y_pred})
predict.head()
```

```
Out[84]:
```

	Actual	Predicted
0	0	0
1	1	1
2	1	1
3	0	1
4	1	1

```
In [88]: # Confusion Matrix to check the Model accuracy
from sklearn.metrics import confusion_matrix
cm = confusion_matrix (y,y_pred)
cm
```

```
Out[88]: array([[435, 250],
               [147, 508]], dtype=int64)
```

```
In [91]: # Calculating Accuracy for the model
# Accuracy = (TP + TN / TP + TN + FP +FN )* 100
((435+508)/(435+250+147+508))*100
```

```
Out[91]: 70.3731343283582
```

```
In [94]: # Classification Report
from sklearn.metrics import classification_report
print(classification_report(y,y_pred))
```

	precision	recall	f1-score	support
0	0.75	0.64	0.69	685
1	0.67	0.78	0.72	655
accuracy			0.70	1340
macro avg	0.71	0.71	0.70	1340
weighted avg	0.71	0.70	0.70	1340

```
In [ ]: # ROC curve
```


```
In [96]: from sklearn.metrics import roc_curve, roc_auc_score
```

```
In [105... fpr, tpr, thresholds = roc_curve(y, classifier.predict_proba(x)[: ,1])

auc = roc_auc_score(y, y_pred)

plt.plot(fpr, tpr, color='red', label='logistic model (area = %0.2f)'%auc)
plt.plot([0,1], [0,1], 'k--')
plt.xlabel('False Positive Rate or [1- True Positive Rate]')
plt.ylabel('True Positive Rate')
```

Out[105...



In [109...

```
classifier.predict_proba(x)[: ,0]
```

Out[109...


```
array([0.99997503, 0.4913874 , 0.42154373, ..., 0.28153744, 0.68381839,  
       0.33466588])
```

In [106...

```
fpr, tpr, thresholds = roc_curve(y, classifier.predict_proba(x)[: ,0])  
  
auc = roc_auc_score(y, y_pred)  
  
plt.plot(fpr, tpr, color='red', label='logistic model (area = %0.2f)'%auc)  
plt.plot([0,1], [0,1], 'k--')  
plt.xlabel('False Positive Rate or [1- True Positive Rate]')  
plt.ylabel('True Positive Rate')
```

Out[106...

```
Text(0, 0.5, 'True Positive Rate')
```



In [107...

```
auc
```

Out[107...

```
0.7053045077171672
```

In [108...

```
classifier.predict_proba(x)
```

Out[108...

```
array([[9.99975026e-01, 2.49742515e-05],  
       [4.91387400e-01, 5.08612600e-01],  
       [4.21543727e-01, 5.78456273e-01],  
       ...,  
       [2.81537440e-01, 6.83818390e-01],  
       [3.34665880e-01, 6.65334666e-01],  
       [9.99997503e-01, 2.49975026e-05]])
```

```
...  
[2.81537440e-01, 7.18462560e-01],  
[6.83818393e-01, 3.16181607e-01],  
[3.34665878e-01, 6.65334122e-01]])
```

```
In [111... y_pred.reshape(-1,1)
```

```
Out[111... array([[0],  
[1],  
[1],  
...  
[1],  
[0],  
[1]], dtype=int64)
```

```
In [ ]:
```

```
In [ ]:
```